

Nomad Notes

- 1 Overview
- 2 Installation
- 3 Networking Information
- 4 Nomad Agent Daemon Configuration
 - 4.1 Nomad Server v0.5.2
 - 4.2 Nomad Client v0.5.2
- 5 Anatomy of a Nomad Job Specification
- 6 Usage
 - 6.1 Display version
 - 6.2 Start up nomad with a config
 - 6.3 Get status of a node running as a client
 - 6.4 Show server-members
 - 6.5 Testing out nomad from the cli
 - 6.6 Validate a nomad job
 - 6.7 Plan and diff a job
 - 6.8 Show previously run jobs and their status
- 7 Troubleshooting
 - 7.1 Resource pool members "Down"
- 8 References

Overview



Nomad is a cluster manager, designed for both long lived services and short lived batch processing workloads. Developers use a declarative job specification to submit work, and Nomad ensures constraints are satisfied and resource utilization is optimized by efficient task packing. Nomad supports all major operating systems and virtualized, containerized, or standalone applications.

Installation

<https://www.nomadproject.io/downloads.html>

```
wget
https://releases.hashicorp.com/nomad/0.5.2/nomad_0.5.2_linux_amd64.zip
unzip nomad_0.5.2_linux_amd64.zip
rm -f nomad_0.5.2_linux_amd64.zip
mv nomad /usr/bin/nomad
mkdir -p /etc/nomad
mkdir -p /opt/nomad
nomad agent -config /etc/nomad/server.hcl
```

Networking Information

Protocol	Port Number	Description	Role that uses this port
TCP	4646	The port used for internal RPC communication between agents and servers, and for inter-server traffic for the consensus algorithm (raft).	Server
TCP	4647	The port used to run the HTTP server. The command line client interacts with this to retrieve data.	Client, Server
TCP	4648	The port used for the SERF gossip protocol for cluster membership. Both TCP and UDP should be routable between the server nodes on this port.	Server
UDP	4648	The port used for the SERF gossip protocol for cluster membership. Both TCP and UDP should be routable between the server nodes on this port.	Server

Nomad Agent Daemon Configuration

Nomad Server v0.5.2

```

log_level    = "INFO"
data_dir     = "/opt/nomad/"
region       = "us-east-1"
bind_addr    = "192.168.33.10"
datacenter   = "molecule"
name         = "centos-7-srv1"

# Takes precedence over bind_addr
addresses {
  http = "192.168.33.10"
  rpc  = "192.168.33.10"
  serf = "192.168.33.10"
}

# Takes precedence over bind_addr
advertise {
  # We need to specify our host's IP because we can't
  # advertise 0.0.0.0 to other nodes in our cluster.
  http = "192.168.33.10"
  rpc  = "192.168.33.10"
  serf = "192.168.33.10"
}

server {
  enabled           = true
  bootstrap_expect  = 1
  encrypt           = "mooleecuulmooleecuulmo=="
}

consul {
  address = "192.168.33.10:8500"
}

# Defaults
ports {
  http = "4646"
  rpc  = "4647"
  serf = "4648"
}

```

Nomad Client v0.5.2

<https://www.nomadproject.io/docs/agent/configuration/client.html>

```

log_level    = "INFO"
data_dir     = "/opt/nomad/"
region       = "us-east-1"
datacenter   = "molecule"
name         = "centos-7-cl1"

# Takes precedence over bind_addr
addresses {
  http = "192.168.33.13"
  rpc   = "192.168.33.13"
  serf  = "192.168.33.13"
}

client {
  enabled = true
  servers = ["192.168.33.10"]

  reserved {
    cpu           = 512
    memory        = 768
    disk          = 10240
    reserved_ports = "22,80,443,4646-4648,8500-8600"
  }
}

consul {
  address = "192.168.33.13:8500"
}

# Defaults
ports {
  http = "4646"
  rpc   = "4647"
  serf  = "4648"
}

```

Anatomy of a Nomad Job Specification

<https://www.nomadproject.io/docs/job-specification/service.html>

job -> group -> task -> service

Key	Description
Job	The <code>job</code> stanza is the top-most configuration option in the job specification. A job is a declarative specification of tasks that Nomad should run. Jobs have a globally unique name, one or many task groups, which are themselves collections of one or many tasks.
Group	The <code>group</code> stanza defines a series of tasks that should be co-located on the same Nomad client. Any <code>task</code> within a group will be placed on the same client.

Task	The <code>task</code> stanza creates an individual unit of work, such as a Docker container, web application, or batch processing.
Service	The <code>service</code> stanza instructs Nomad to register the task as a service using the service discovery integration.

Usage

Display version

```
nomad --version
```

Start up nomad with a config

The config determines whether or not the node acts as a client or server despite the command being **nomad agent**

```
nomad agent -config /etc/nomad/server.hcl
```

Get status of a node running as a client

```
nomad node-status -self
```

From a client

```
[root@ip-252-25-69-190 nomad]# nomad node-status -self
ID      = 74c4ef9c
Name    = ip-252-25-69-190.internal
Class   = <none>
DC      = dc1
Drain    = false
Status  = ready
Uptime  = 2158h35m31s
```

Allocated Resources

CPU	Memory	Disk	IOPS
0/4988 MHz	0 B/3.9 GiB	0 B/38 GiB	0/0

Allocation Resource Utilization

CPU	Memory
0/4988 MHz	0 B/3.9 GiB

Host Resource Utilization

CPU	Memory	Disk
0/4988 MHz	1.2 GiB/3.9 GiB	12 GiB/50 GiB

From a server

```
[root@ip-252-25-5-53 nomad]# nomad node-status
-address=http://172.25.5.53:4646
ID      DC      Name                                     Class
Drain    Status
74c4ef9c dc1    ip-252-25-69-190.internal    <none>  false  ready

[root@centos-7-srv1 ~]# nomad node-status
ID      DC      Name          Class  Drain  Status
96df05b8 molecule centos-7-cl1  <none> false  ready
```

Show server-members

```
nomad server-members -address=http://172.25.5.53:4646
nomad server-members
```

```
[root@ip-252-25-5-53 nomad]# nomad server-members
-address=http://172.25.5.53:4646
Name                                     Address
Port  Status  Leader  Protocol  Build  Datacenter  Region
ip-252-25-5-53.internal.us-east-1  172.25.5.53  4648  alive   true    2
0.4.1  dc1          us-east-1

[root@centos-7-cl1 ~]# nomad server-members
Name                                     Address          Port  Status  Leader  Protocol
Build  Datacenter  Region
centos-7-srv1.us-east-1  192.168.33.10  4648  alive   true    2
0.5.1  molecule     us-east-1
```

Testing out nomad from the cli

```
nohup nomad agent -config /etc/nomad/server.hcl &>nomad.log &
tailf nomad.log
```

Validate a nomad job

You should check out the anatomy of a nomad job! <https://www.nomadproject.io/docs/jobspec/>

```
nomad validate $JOBNAME.nomad
```

```
[root@ip-252-25-5-53 ~]# nomad validate sleep.nomad
Job validation successful
```

This particular job file looks like

```

job "sleepy-job" {
  region = "us-east-1"
  datacenters = ["dc1"]
  type = "batch"

  constraint {
    attribute = "${attr.kernel.name}"
    value = "linux"
  }

  task "gotosleep" {
    driver = "exec"
    config {
      command = "/bin/sleep"
      args = ["1"]
    }
    resources {
      cpu = 200
    }
  }
}

```

On a receiving client, the executing job from above will look as follows. Pay attention to the task name **gotosleep**

```

[root@ip-252-25-69-190 ~]# ps aux | egrep '(sleep|nomad)'
root      12250 59.7  0.5 390768 22164 ?        Ssl  21:07   0:02
/usr/bin/nomad executor
/opt/nomad/alloc/40b04879-5eb0-62b8-088d-fdfb8cdb3b4d/gotosleep/gotosleep-
executor.out
nobody    12273  0.0   0.0 107908   676 ?        S    21:07   0:00 bin/sleep
20

```

Plan and diff a job

This is a very similar concept to Terraform. You'll want to do this prior to running the actual job.

```

[root@ip-252-25-5-53 ~]# nomad plan -diff
-address=http://172.25.5.53:4646 sleep.nomad

```



```
[root@ip-252-25-5-53 ~]# nomad plan -diff -address=http://172.25.5.53:4646 sleep.nomad
```

```
[root@centos-7-srv1 ~]# nomad plan example.nomad
```

```
+ Job: "sleepy-job"
```

```
+ Task Group: "gotosleep" (1 create)
```

```
  + Task: "gotosleep" (forces create)
```

Scheduler dry-run:

- All tasks successfully allocated.

Job Modify Index: 0

To submit the job with version verification run:

```
nomad run -check-index 0 sleep.nomad
```

When running the job with the check-index flag, the job will only be run if the

server side version matches the job modify index returned. If the index has changed, another user has modified the job and the plan's results are potentially invalid.

Show previously run jobs and their status

```
nomad status
```

```
nomad status $JOBID
```

```
nomad inspect $JOBID
```

```
[root@ip-252-25-5-53 ~]# nomad status
ID           Type     Priority  Status
sleepy-job   batch   50       dead
writer       batch   50       dead

[root@ip-252-25-5-53 ~]# nomad status writer
ID           = writer
Name         = writer
Type         = batch
Priority      = 50
Datacenters  = dc1
Status       = dead
Periodic     = false

Summary
Task Group  Queued  Starting  Running  Failed  Complete  Lost
echo        0        0          0         1         1         0

Allocations
ID           Eval ID   Node ID   Task Group  Desired  Status    Created At
2e7cd21f    f60c1d1f  74c4ef9c  echo        run      failed    10/18/16
21:57:01 EDT
188def87    742d43aa  74c4ef9c  echo        stop     complete  10/18/16
21:54:28 EDT
[root@ip-252-25-5-53 ~]# nomad status sleepy-job
ID           = sleepy-job
Name         = sleepy-job
Type         = batch
Priority      = 50
Datacenters  = dc1
Status       = dead
Periodic     = false

Summary
Task Group  Queued  Starting  Running  Failed  Complete  Lost
gotosleep   0        0          0         0         3         0

Allocations
ID           Eval ID   Node ID   Task Group  Desired  Status    Created At
40b04879    e5ccc8c8  74c4ef9c  gotosleep   run      complete  10/18/16
21:07:40 EDT
15752502    ff78cb88  74c4ef9c  gotosleep   stop     complete  10/18/16
21:06:41 EDT
d31db0b0    ecc0aa67  74c4ef9c  gotosleep   stop     complete  10/18/16
21:05:29 EDT
```

Troubleshooting

Resource pool members "Down"

If you are seeing the following message on a client

```
$ cat /var/log/nomad/nomad.log
==> Error starting agent: client setup failed: failed to restore state: 2
error(s) occurred:

* failed to decode state: unexpected end of JSON input
* failed to decode state: unexpected end of JSON input
  Loaded configuration from /etc/nomad/client.hcl
==> Starting Nomad agent...
==> Error starting agent: client setup failed: failed to restore state: 2
error(s) occurred:

* failed to decode state: unexpected end of JSON input
* failed to decode state: unexpected end of JSON input
  Loaded configuration from /etc/nomad/client.hcl
==> Starting Nomad agent...
```

and you are seeing the following message on a server

```
[root@centos-7-srv1 jobs]# nomad node-status
ID            DC          Name           Class   Drain  Status
73cf3fc7      molecule    centos-7-cl3   <none>  false  down
e3496454      molecule    centos-7-cl1   <none>  false  down
a42fd16a      molecule    centos-7-cl2   <none>  false  ready
```

You should then run the following commands on the affected clients

```
systemctl stop nomad
rm -rf /opt/nomad/client/alloc/*
rm -rf /opt/nomad/alloc/*
systemctl restart nomad
```

and watch the server for the updated node-status

```
[root@centos-7-srv1 jobs]# nomad node-status
```

ID	DC	Name	Class	Drain	Status
73cf3fc7	molecule	centos-7-cl3	<none>	false	down
e3496454	molecule	centos-7-cl1	<none>	false	down
a42fd16a	molecule	centos-7-cl2	<none>	false	ready


```
[root@centos-7-srv1 jobs]# nomad node-status
```

ID	DC	Name	Class	Drain	Status
73cf3fc7	molecule	centos-7-cl3	<none>	false	initializing
e3496454	molecule	centos-7-cl1	<none>	false	initializing
a42fd16a	molecule	centos-7-cl2	<none>	false	ready


```
[root@centos-7-srv1 jobs]# nomad node-status
```

ID	DC	Name	Class	Drain	Status
73cf3fc7	molecule	centos-7-cl3	<none>	false	ready
e3496454	molecule	centos-7-cl1	<none>	false	ready
a42fd16a	molecule	centos-7-cl2	<none>	false	ready

References

<https://github.com/hashicorp/nomad>

<https://www.nomadproject.io/>

<https://www.nomadproject.io/docs/job-specification/service.html>